

LECTURE 1.3

CPU SCHEDULING

ALGORITHMS I

COP4600

Dr. Matthew Gerber

2/1/2016

First-Come First-Served (6.3.1)

- The simplest of all reasonable schedulers
- Processes are allocated the CPU in the order in which they arrive
- Processes run until completion and termination
- Cannot be sensibly made preemptive
- Strengths
 - Really simple
 - Easy to write, easy to understand
- Weaknesses
 - Everything else
 - Completely unpredictable wait times
 - Non-preemptive nature makes it unsuitable for use in modern OSs

First-Come First-Served (6.3.1)

```
define pcb {  
    int                process_id  
    enumeration        state  
    register           program_counter  
    list<register>      registers  
    scb                schedule  
    mcb                memory  
    acb                account  
    iocb               io  
}
```

```
define scb {  
    time              arrival  
}
```

```
queue<pcb> ready_queue;
```

```
function proc_ready(pcb &p) {  
    p.state = READY  
    p.schedule.arrival = now()  
    ready_queue.enqueue(p)  
}
```

```
function select_proc() {  
    pcb p  
  
    p = ready_queue.dequeue()  
    dispatch(p)  
}
```

First-Come First-Served (6.3.1)

```
define pcb {  
    int            process_id  
    enumeration    state  
    register       program_counter  
    list<register> registers  
    scb            schedule  
    mcb            memory  
    acb            account  
    iocb           io  
}
```

```
define scb {  
    time           arrival  
}
```

```
list<pcb> ready_queue;
```

```
function proc_ready(pcb &p) {  
    p.state = READY  
    p.schedule.arrival = now()  
    ready_queue.add(p)  
}
```

```
function select_proc() {  
    pcb p  
  
    p = find_minimum<pcb>(  
        each rqp in ready_queue,  
        rqp.schedule.arrival  
    )  
    ready_queue.delete(p)  
  
    dispatch(p)  
}
```

First-Come First-Served (6.3.1)

Process	Burst Time	Wait Time
P1	24	0
P2	3	24
P3	3	27



First-Come First-Served (6.3.1)

Process	Burst Time	Wait Time
P1	3	0
P2	3	3
P3	24	6



Shortest Job First (6.3.2)

- Processes are scheduled to the CPU giving priority to the process that will have the next-shortest burst
- Can be left cooperative or made preemptive
- Strengths
 - Guaranteed to produce an optimal schedule in terms of waiting time
- Weaknesses
 - Not actually possible
 - Longer processes may never get run (we'll look at this more in a bit...)

Shortest Job First (6.3.2)

```
define pcb {  
    int                process_id  
    enumeration        state  
    register           program_counter  
    list<register>      registers  
    scb                schedule  
    mcb                memory  
    acb                account  
    iocb               io  
}
```

```
define scb {  
    time              est_length  
}
```

```
list<pcb> ready_queue;
```

```
function proc_ready(pcb &p) {  
    p.state = READY  
    time t = estimate_time(p)  
    p.schedule.est_length = t  
    ready_queue.add(p)  
}
```

```
function select_proc() {  
    pcb p  
  
    p = find_minimum<pcb>(  
        each rqp in ready_queue,  
        rqp.schedule.est_length  
    )  
    ready_queue.delete(p)  
  
    dispatch(p)  
}
```


Shortest Job First (6.3.2) (Non-preemptive)

Process	Burst Time	Wait Time
P1	6	3
P2	12	18
P3	9	9
P4	3	0



Shortest Job First (6.3.2) (Preemptive)

Process	Burst Time	Arrival	Wait Time
P1	6	0	3
P2	12	1	17
P3	9	2	7
P4	3	3	0



Priority (6.3.3)

- More general version of Shortest Job First
 - (Actually, it's the other way around...)
- Each process has a *priority* that determines how soon it is scheduled compared to other processes
- This priority can either be a number or determined by some sort of function
- Again, can be left cooperative or made preemptive
- Strengths
 - Flexible
- Weaknesses
 - Completely dependent on the priority being set reasonably
 - What happens if some process has a priority bad enough that some other process always has better?
 - This effect is called *starvation*

Priority (6.3.3)

```
define pcb {  
    int                process_id  
    enumeration        state  
    register           program_counter  
    list<register>      registers  
    scb                schedule  
    mcb                memory  
    acb                account  
    iocb               io  
}
```

```
define scb {  
    int                priority  
}
```

```
list<pcb> ready_queue;
```

```
function proc_ready(pcb &p) {  
    p.state = READY  
    ready_queue.add(p)  
}
```

```
function select_proc() {  
    pcb p  
  
    p = find_minimum<pcb>(  
        each rqp in ready_queue,  
        rqp.schedule.priority  
    )  
    ready_queue.delete(p)
```

```
    dispatch(p)  
}
```

Priority (6.3.3) (Non-preemptive)

Process	Burst Time	Priority	Wait Time
P1	6	2	12
P2	12	1	0
P3	9	4	21
P4	3	3	18



Round-Robin (6.3.4)

- Algorithm explicitly designed for time-sharing systems
- Some similarities to First Come First Served
- Redesigned fundamentally to add preemption
- Define a *time quantum* or *time slice* (generally between 10ms and 100ms)
- Set the timer interrupt to stop the running process if its burst goes longer than the time quantum
- When a process becomes ready for whatever reason, put it on the back of the queue

Round-Robin (6.3.4)

```
define pcb {  
    int                process_id  
    enumeration        state  
    register           program_counter  
    list<register>      registers  
    scb                schedule  
    mcb                memory  
    acb                account  
    iocb               io  
}
```

```
define scb {  
    time              arrival  
}
```

```
queue<pcb> ready_queue;
```

```
function proc_ready(pcb &p) {  
    p.state = READY  
    p.schedule.arrival = now()  
    ready_queue.enqueue(p)  
}
```

```
function timer_int(pcb &current) {  
    proc_ready(current)  
    select_proc()  
}
```

```
function select_proc() {  
    pcb p  
  
    p = ready_queue.dequeue()  
    dispatch(p)  
}
```

Round Robin (6.3.4) ($q = 2$)

Process	Burst Time	Wait Time
P1	13	7
P2	4	7
P3	2	4
P4	1	6

P1		P2		P3		P4	P1		P2		P1		P1		P1		P1		P1	
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20

Round Robin (6.3.4)

- The basis of the scheduling algorithm for most interactive time-sharing systems
 - The *basis*, not what they actually use
- Not too great on overall wait time, but...
- Each process is guaranteed to have to wait at most $(n-1)q$ time units until it gets the processor back
- q needs to be large with respect to the context switch time
 - But that said, in the modern context, context switching *actually* takes a few *microseconds*
- q needs to be small enough for the system to be continually responsive
 - Round Robin degenerates to First Come First Served with a large enough q

NEXT TIME: SCHEDULING ALGORITHMS II
