

LECTURE 1.5

SCHEDULING CASE STUDIES AND INTRODUCTION TO IPC

COP4600

Dr. Matthew Gerber

2/8/2016

SCHEDULING CASE STUDIES

Windows Scheduling (6.7.2)

- Explicit multilevel queue with internal priorities mapped to linear numeric values that in turn provide an implicit multilevel queue
 - It's probably better to just show what we mean; we'll do that next
- In general, higher-priority threads immediately preempt lower-priority threads
- In particular, (soft) real-time threads *always* preempt other threads

Windows Scheduling (6.7.2)

Priority Classes

- REALTIME_PRIORITY_CLASS
- HIGH_PRIORITY_CLASS
- ABOVE_NORMAL_PRIORITY_CLASS
- NORMAL_PRIORITY_CLASS
- BELOW_NORMAL_PRIORITY_CLASS
- IDLE_PRIORITY_CLASS

Relative Thread Priorities

- TIME_CRITICAL
- HIGHEST
- ABOVE_NORMAL
- NORMAL
- BELOW_NORMAL
- LOWEST
- IDLE

Windows Scheduling (6.7.2)

	REAL TIME	HIGH	ABOVE_ NORMAL	NORMAL	BELOW_ NORMAL	IDLE
TIME_ CRITICAL	31	15	15	15	15	15
HIGHEST	26	15	12	10	8	6
ABOVE_ NORMAL	25	14	11	9	7	5
NORMAL	24	13	10	8	6	4
BELOW_ NORMAL	23	12	9	7	5	3
LOWEST	22	11	8	6	4	2
IDLE	16	1	1	1	1	1

Windows Scheduling (6.7.2)

- When a thread returns from an I/O wait, its priority is raised
 - High amount for user interaction
 - Lower amount for disk I/O
- When a thread uses its entire quantum, its priority is lowered
 - Never lowered below base priority
- The process owning the foreground window gets a priority bump
 - Usually about 3

Linux Scheduling: $O(1)$

- Linux scheduler for kernel version 2.6 up until 2.6.23
- The bad: Heuristics were pretty bad at figuring out which processes were interactive
- The good: An $O(1)$ scheduler!
- Any priority queue “should” require $O(\log n)$, at least, and the previous Linux scheduler was $O(n)$
- So this scheduler implements priority with a guaranteed constant upper bound of execution time
 - How?
- A constant number of priority levels, *and a queue for each of them*

Linux Scheduling: $O(1)$

- 140 priority levels
- Each processor is given two arrays of 140 doubly linked lists: the *active* and *expired* array
 - When a process's quantum expires it's re-inserted into the corresponding queue in the expired array
 - Its priority can change at this point
 - When the active array runs out of processes, the scheduler *switches arrays*
- When a process needs to be chosen, an active process with best priority is chosen
- Guaranteed to be doable in 140 compares or less (not counting dereferencing)...
 - So this is technically $O(1)$
 - Very, very, very technically

Linux Scheduling: CFS (6.7.1)

- The Completely Fair Scheduler
- Implements the *weighted fair queuing* concept
 - Given equal priorities, processes should receive as close to the same amount of time from the processor as possible
- Each process is given a *virtual runtime*
- The scheduler selects whichever process currently has the *lowest* virtual runtime
- For processes with completely normal priority, virtual runtime is equal to physical runtime

- Virtual runtime is affected by multipliers
 - Priority is a direct multiplier
 - Long-running processes get de-prioritized
 - CPU-bound processes get de-prioritized
 - I/O bound processes get prioritized
- This *is* all done preemptively
- Guards are put in place (maximum time a process can wait to run, minimum time a process can run) to prevent silly results

Linux Scheduling: CFS (6.7.1)

- [illegible]

Solaris Scheduling (6.7.3)

- Solaris traditionally has four classes of thread:
 - Real-Time
 - System
 - Time-Sharing
 - Interactive
- Solaris 9 adds two more:
 - Fair Share
 - Fixed Priority
- *All of these use different scheduling methods*
- The scheduler maps each class-specific output to a global priority
- Real-Time threads always have priority over everything else
- System threads have priority over everything except Real-Time threads

Solaris Scheduling (6.7.3)

- Interactive/Time-Sharing schedules are similar, and are the most interesting for our purposes
- Solaris uses a dynamic priority method somewhat similar to Windows'
 - (Actually, it's the other way around)
- Threads that return from I/O have their priority increased (in Solaris' case, radically)
- Threads that use up their time quanta have their priority decreased (more gradually)
- Threads with lower priority get longer time quanta

INTER-PROCESS COMMUNICATION

Inter-Process Communication (3.5)

- Two fundamental mechanisms available:
 - Shared Memory
 - Message Passing
- Both of these are useful, and many operating systems provide both
- Message-passing avoids conflicts and works better across networks
- Shared memory is more convenient for large amounts of data
- But we'll get back to this...

The Producer-Consumer Problem

- Classic example of inter-process sharing
- A *producer* in one process creates chunks of data for a *consumer* to further operate on
- The producer needs to be able to place those chunks of data in a shared memory location that the consumer can read them from
- The producer needs to be delayed as appropriate to not overrun this buffer
- The consumer needs to be delayed as appropriate to wait for the producer to generate data for it to operate on
- So how do we make this happen?

The Producer-Consumer Problem

- Classic example of inter-process sharing
- A *producer* in one process creates chunks of data for a *consumer* to further operate on
- The producer needs to be able to place those chunks of data in a shared memory location that the consumer can read them from
- The producer needs to be delayed as appropriate to not overrun this buffer
- The consumer needs to be delayed as appropriate to wait for the producer to generate data for it to operate on
- So how do we make this happen?

Producer-Consumer Code

```
constant BUFSIZE = 8
shared data buffer[BUFSIZE]
shared int in = 0
shared int out = 0
```

```
producer_thread {
    data d
    while true {
        while ((in+1)%BUFSIZE) == out {
            wait
        } // while in + 1 == out

        d = produce_item()
        buffer[out] = d
        in = in + 1
        in = in % BUFSIZE
    } // while true
}
```

```
consumer_thread {
    data d
    while true {
        while in == out {
            wait
        } // while in == out

        data d = buffer[out]
        consume_item(d)
        out = out + 1
        out = out % BUFSIZE
    } // while true
}
```

Producer-Consumer Code

```
constant BUFSIZE = 8
shared data buffer[BUFSIZE]
shared int in = 0
shared int out = 0
```

```
producer_thread {
    data d
    while true {
        while ((in+1)%BUFSIZE) == out {
            wait
        } // while in + 1 == out

        d = produce_item()
        buffer[out] = d
        in = in + 1
        in = in % BUFSIZE
    } // while true
}
```

Do you see the problem?

```
consumer_thread {
    data d
    while true {
        while in == out {
            wait
        } // while in == out

        data d = buffer[out]
        consume_item(d)
        out = out + 1
        out = out % BUFSIZE
    } // while true
}
```

Producer-Consumer Code

```
constant BUFSIZE = 8
shared data buffer[BUFSIZE]
shared int in = 0
shared int out = 0
```

```
producer_thread {
    data d
    while true {
        while ((in+1)%BUFSIZE) == out {
            wait
        } // while in + 1 == out

        d = produce_item()
        buffer[out] = d
        in = in + 1
        in = in % BUFSIZE
    } // while true
}
```

**Here's a
hint...**

```
consumer_thread {
    data d
    while true {
        while in == out {
            wait
        } // while in == out

        data d = buffer[out]
        consume_item(d)
        out = out + 1
        out = out % BUFSIZE
    } // while true
}
```

Producer-Consumer Code

```
constant BUFSIZE = 8
shared data buffer[BUFSIZE]
shared int in = 0
shared int out = 0
```

```
producer_thread {
  data d
  while true {
    while ((in+1)%BUFSIZE) == out {
      wait
    } // while in + 1 == out

    d = produce_item()
    buffer[out] = d
    in = in + 1
    in = in % BUFSIZE
  } // while true
}
```

**What happens if we
context-switch inside
one of these sections?**

```
consumer_thread {
  data d
  while true {
    while in == out {
      wait
    } // while in == out

    data d = buffer[out]
    consume_item(d)
    out = out + 1
    out = out % BUFSIZE
  } // while true
}
```

Producer-Consumer Code

```
constant BUFSIZE = 8
shared data buffer[BUFSIZE]
shared int in = 0
```

s
p

BAD THINGS

```
wait
} // while in + 1 == out
```

```
d = produce_item()
buffer[out] = d
in = in + 1
in = in % BUFSIZE
} // while true
}
```

```
wait
} // while in == out
```

```
data d = buffer[out]
consume_item(d)
out = out + 1
out = out % BUFSIZE
} // while true
}
```

The Critical Section Problem

- This kind of section of code actually has a name
- It's called a *critical section*
- The hard part of synchronization is keeping critical sections from interfering with each other

NEXT TIME:
SYNCHRONIZATION
