

INTRODUCTION 2

MEMORY, I/O, AND OPERATING ENVIRONMENTS

COP4600

Dr. Matthew Gerber

1/20/2016

THE MODERN OPERATING SYSTEM, PART II

Memory Management (1.7)

- Main memory is the largest storage device the CPU can directly address
 - There are exceptions here, but all of them are special cases
- Programs have to be in main memory to run
- The operating system must...
 - Keep track of what memory is being used, and by which processes
 - Decide what to move into and out of memory
 - Allocate and de-allocate memory as necessary
 - Make sure processes don't "color outside the lines" of their memory (more on this later)

Input and Output: Files and Disks (1.8)

- A file is *a named, sequential array of bytes*.
- Almost all data on disks are files
 - There *are* exceptions
 - They're usually esoteric
- The operating system has to manage...
 - Creating, writing, reading and deleting files
 - Creating, writing, reading and deleting directories to organize files
 - Providing functions for processes to work with files
 - Mapping files onto disks
 - (really, partitions, but we'll call them disks for now)

Input and Output: Files and Disks (1.8)

- The operating system *also* has to manage disks themselves:
 - Organizing files on disks
 - Managing free space on disks
 - Allocating space on disks to files
 - Scheduling disks for processes
- By the way, technically a disk is just another I/O device
- They're so foundational we tend to think of them separately

Input and Output: Everything Else

- *Almost* all other devices in a modern operating system are viewed in one of two ways:
 - As a memory space
 - As a file
- ***Can you think of the biggest exception?***
 - (and even it's not totally immune...)

Input and Output: Everything Else

- Each device has a *device driver* that either:
 - Maps it to a file
 - Maps it to memory
 - Maps it to another core OS service (like **networking**)
 - Gives it a unique interface (very rare in a modern OS)

Security (1.9, sort of): Authentication

- *Most* modern operating systems allow more than one user
 - There are prominent exceptions to this
 - Allowing more than one user *at a time* is *not* ubiquitous yet
- Even operating systems that only allow one user usually have the concept of a user identity
- *Authentication* is the process of determining that *you are who you say you are*.
- Allowing users to misidentify themselves leads to all sorts of calamity, whether accidental or malicious

Security (1.9 – sort of): Authentication

- The classic authentication method is the password – something you *know*
 - Less useful than it sounds
 - Good passwords are hard to remember
 - (And by the contrapositive...)
 - Passphrases are (much) better but are not a panacea

Security (1.9 – sort of): Authentication

- Token-based, or classic *two-factor* authentication, adds something you *have*
 - When properly implemented, makes authentication much harder to defeat
 - *Much harder is not the same thing as impossible*
 - “Tokens” that are really computers can be compromised themselves
 - Be smart about what you allow to authenticate what

Security (1.9 – sort of): Authentication

- Biometric authentication adds something you *are*
 - Very difficult to bypass without personalized effort
 - Also very difficult to reliably implement

Security (1.9 – sort of): Authorization

- Once you know who somebody is you still need to know what that means
- *Authorization* is the process of determining, given who you are, *what you should be able to do*
- Certain users are able to perform certain actions, and not others
 - Classic example: “You can only access your own files”
- This needs to be correctly determined *and* successfully enforced for the system to remain secure

Security (1.9 – sort of): Authorization

- Certain users may have the right to perform other actions if they *escalate privileges* first
 - Classic example: UNIX-like systems' "Root access"
 - One method is to actually log in to the root account
 - Another is to explicitly escalate to root privileges temporarily, such as with **sudo**
 - Windows version: "Administrator access"
 - Older versions of Windows allow administrator actions immediately
 - User Account Control (UAC) in newer versions tries to put some brakes on
 - Other systems handle in different ways
 - The text uses **setuid** as an example...

Security (1.9 – sort of): Authorization

- Modern authorization makes heavy use of access control lists and grouping
- Authorization *seems* less dangerous than authentication...
- ...but a lot of system-level attacks take place against badly designed authorization systems
 - Ever heard of a “privilege-escalation attack”?

Security (1.9 – sort of): Cryptography

- Practical cryptography is intended to do two things:
 - Allow you to send messages that are very hard for anybody other than your recipient to read
 - Allow you to sign messages in a way that is very hard for anybody else to impersonate
- Operating systems use cryptography for, among other things:
 - Secure storage in the face of potential theft
 - Secure communication over potentially hostile networks
- **Be realistic about the limits of cryptography!**

Putting It Together

- Modern operating systems need:
 - Process Management
 - Memory Management
 - I/O Management
 - Security
- ...and how many places does all this show up?

OPERATING ENVIRONMENTS (1.11)

Physical Computers

- Large Systems
 - Initially required to centralize very scarce computing resources
 - Built around time-sharing; job-oriented more than interactivity-oriented
 - Decreasingly common, but not dead yet
- Desktop/Laptop Computing
 - What most people think of as “operating systems” at this point
 - Interrupt-driven, multitasking, etc.
 - Usually supports multiple users, may or may not support multiple at once

Physical Computers

- Mobile Computing
 - Laptops and tablets
 - Have a lot of the characteristics of desktop operating systems
 - Tend to be optimized for information consumption
- Real-Time Operating Systems
 - When things have to be done *now*
 - Popular for control systems
 - You really don't want your X-Ray machine to decide to garbage collect at random times
 - Can (and do) show up as subsets of other systems

Virtual Computers

- *Emulation* is running code for another kind of computer
- *Interpretation* is running code for a theoretical computer
 - To note, calling Java “interpreted” is a stretch
- *Virtualization* is running a computer on a computer
 - Virtualization generally requires specialized hardware support...
 - ...but almost all modern desktop CPUs have it

Networked Computers

- Client-Server Systems
 - Networked systems that allow a *client* computer to ask a *server* to do something on its behalf
 - Easy example: Web servers
 - More complicated example: Database servers
 - Still more complicated example: Game servers

Networked Computers

- Peer-to-Peer Systems
 - Actually a lot like client-server computing, except at any time, any computer can be a client, a server or both
 - Computers have to be able to discover each other, either via a central directory or exchanging information over the network
 - Famous for piracy, but there are plenty of other applications
 - There *are* reasons other than piracy that network administrators don't like it

Networked Computers

- Cloud Computing
 - Broadly constructed client-server systems running over the commodity Internet or private networks
 - Can include file servers, computation servers, virtual machines, etc.
 - Not a new technology so much as a radically scaled application of existing technologies
 - The innovation is mostly in...
 - ...the scale itself, and
 - ...the ability to usefully control client-server computing *at* that scale

NEXT TIME:
PROCESSES AND THREADS
