

LECTURE 1.6

SYNCHRONIZATION

COP4600

Dr. Matthew Gerber

2/10/2015

Review: Producer/Consumer

- Two threads that each use a shared, bounded buffer
- The producer produces data items for the consumer to consume
- The producer needs to avoid overrunning the buffer
 - Not hard; just use shared variables
- The consumer needs to avoid trying to consume items that aren't ready yet
 - Also not hard; see above
- Both need to avoid messing up each others' updating of the aforementioned shared variables
- *That's* the hard part
- Again, this problem has a general name

Critical Sections (5.2)

- *A critical section* is:
 - A *segment of code*...
 - In a *cooperating process*...
 - That *modifies shared information* in a way that *either*...
 - ...*can destructively interfere* with other critical sections, or
 - ...*can be destructively interfered with by* other critical sections
- It is considered (and really only makes sense) in terms of a *given* set of inter-process interactions
- We have three fundamental ways to deal with critical sections
- The first way is really, really simple...

Producer-Consumer with Disabling Interrupts

```
constant BUFSIZE = 8
shared data buffer[BUFSIZE]
shared int in = 0
shared int out = 0
```

```
producer_thread {
    data d
    while true {
        while ((in+1)%BUFSIZE) == out {
            wait
        } // while in + 1 == out

        disable_interrupts()
        d = produce_item()
        buffer[out] = d
        in = in + 1
        in = in % BUFSIZE
        enable_interrupts()
    } // while true
}
```

```
consumer_thread {
    data d
    while true {
        while in == out {
            wait
        } // while in == out

        disable_interrupts()
        data d = buffer[out]
        consume_item(d)
        out = out + 1
        out = out % BUFSIZE
        enable_interrupts()
    } // while true
}
```

Producer-Consumer with Disabling Interrupts

```
constant BUFSIZE = 8
shared data buffer[BUFSIZE]
shared int in = 0
shared int out = 0
```

p

Now with 100% less undefined behavior, but...

```
    disable_interrupts()
    d = produce_item()
    buffer[out] = d
    in = in + 1
    in = in % BUFSIZE
    enable_interrupts()
} // while true
}
```

```
    disable_interrupts()
    data d = buffer[out]
    consume_item(d)
    out = out + 1
    out = out % BUFSIZE
    enable_interrupts()
} // while true
}
```

The Problems with Disabling Interrupts

- Performance
 - If the timer is controlled by interrupts, it's going to mess them up
- Potential unpredictable consequences
 - Disabling fundamental bits of the operating system any more often than absolutely necessary is not a good idea
- Absolute non-starter on multiprocessor systems
 - Consider what it takes to disable interrupts on those...
- So we need a different way to make this happen

The Mutex

- A primitive function provided by the operating system
- The following functions are implemented by the operating system...

```
signal (mutex m) {  
    m = true  
}
```

```
wait (mutex m) {  
    while m == false {  
        wait  
    }  
    m = false  
}
```

- ...in such a way that *wait cannot be interrupted* between *testing* that *m* is true and *setting* it to false (more on this in a bit...)
- As with anything else, a mutex can start in either position

```
mutex m = true
```

Producer-Consumer with a Mutex

```
constant BUFSIZE = 8
shared data buffer[BUFSIZE]
shared int in = 0
shared int out = 0
mutex m = true
```

```
producer_thread {
    data d
    while true {
        while ((in+1)%BUFSIZE) == out {
            wait
        } // while in + 1 == out

        wait(m)
        d = produce_item()
        buffer[out] = d
        in = in + 1
        in = in % BUFSIZE
        signal(m)
    } // while true
}
```

```
consumer_thread {
    data d
    while true {
        while in == out {
            wait
        } // while in == out

        wait(m)
        data d = buffer[out]
        consume_item(d)
        out = out + 1
        out = out % BUFSIZE
        signal(m)
    } // while true
}
```


Producer-Consumer with a Mutex

```
constant BUFSIZE = 8
shared data buffer[BUFSIZE]
shared int in = 0
shared int out = 0
```

m
p

Now with 100% less undefined
behavior *and* system instability

```
} // while in + 1 == out

wait(m)
d = produce_item()
buffer[out] = d
in = in + 1
in = in % BUFSIZE
signal(m)
} // while true
}
```

```
} // while in == out

wait(m)
data d = buffer[out]
consume_item(d)
out = out + 1
out = out % BUFSIZE
signal(m)
} // while true
}
```

Implementing Mutexes (5.5)

- Remember that a mutex basically provides the following functions:

```
signal (mutex m) {  
    m = true  
}
```

```
wait (mutex m) {  
    while m == false {  
        wait  
    }  
    m = false  
}
```

- The only thing that needs to *not be interrupted* is the period in `wait` between determining that `m` is true and setting it to false
- There are a few ways to do this; let's start with the obvious one

Mutex Locks with Disabling Interrupts

```
signal (mutex m) {  
    m = true  
}  
  
wait (mutex m) {  
    disable interrupts  
    while m == false {  
        enable interrupts  
        yield  
        disable interrupts  
    }  
    m = false  
    enable interrupts  
}
```

- `wait` is never interrupted between determining that `m` is true and setting it to false
- In one version of this, `wait` yields; in another version, `wait` just keeps checking whenever it can get the CPU
- Which version to use depends on how long you expect this all to take
- This method – especially the version where you don't yield—is called a *spin wait*
- This all has at least *some* of the problems of disabling interrupts in general

The Test-And-Set instruction (5.4)

- Better way: Get help from the hardware
- The **processor** implements the following instruction **atomically**

```
boolean test_and_set (boolean &target) {  
    boolean rv = target  
    target = true  
  
    return rv  
}
```

- This defines a sort of minimum possible critical section in hardware
- If we call `test_and_set` on a true value, it stays true and we get back true
- If we call `test_and_set` on a false value, it becomes true and we get back false
- It's up to the hardware to make sure no other processors interfere

Mutex Locks with Test-And-Set

```
signal (mutex m) {  
    m = false  
}  
  
wait (mutex m) {  
    while test_and_set(m) {  
        yield  
    }  
}
```

- Note that the logic on this is inverted
- The mutex is free if *m* is *false*; releasing it sets it false, and acquiring its lock sets it true
- When `wait` calls `test_and_set`, it gets back false if *m* was false and true if *m* was true; ***m* is true afterward either way**
- The good news is we're no longer disabling interrupts
- The bad news is we're still yielding or spin-waiting

Mutex Locks: Avoiding Spin-Waits (5.6.2)

- The easy way to solve this is by turning mutexes into queues
- When a process waits on a mutex that isn't available, it's inserted into a queue for that mutex
- When a process signals a mutex that has any processes waiting for it, one is dequeued and given the lock
- The queuing is a critical section in and of itself...
 - ...but we can guard *it* with a spin-wait mutex without any trouble, since it's a very small $O(1)$ problem

Counting Semaphores (5.6)

- Just like mutex locks, except integer instead of boolean
- Signaling increments, waiting decrements; processes are blocked when they wait on zero
- Logic is more complex
 - Requires multiple levels of guarding...
 - ...just like the queued mutex we just discussed
- By the way...
 - How easily could producer-consumer be implemented using counting semaphores?
 - (Hint: Pretty easily)

LOOSE ENDS 1: DEADLOCKS

Deadlocks (5.6.3)

shared mutex m1, m2

```
define p1 {  
    wait(m1)  
    wait(m2)  
    ...
```

```
define p2 {  
    wait(m2)  
    wait(m1)  
    ...
```

- What's going to happen here?

Deadlocks (5.6.3)

shared mutex m1, m2

```
define p1 {  
    wait(m1)  
    wait(m2)  
    ...  
}
```

```
define p2 {  
    wait(m2)  
    wait(m1)  
    ...  
}
```

- What's going to happen here?

BAD THINGS

- A set of processes is *deadlocked* if every process in a set is waiting on an event that can only be caused by another process in the same set
- We'll cover deadlocks in our last lecture of the unit

LOOSE ENDS 2: CASE APPROACHES

Windows Synchronization (5.9.1)

- Uses simple spinlocks and mutex locks in the kernel
- Adapts spinlocks to interrupt-disabling for single-processor systems
- For thread synchronization, Windows provides *dispatcher objects* with:
 - Mutex locks
 - Semaphores
 - Timers
 - Conditions (called Events)
- Lock-guarded shared memory is provided
- *Critical Section Objects* provide mutexes at the user level that only allocate kernel mutexes if they have to

Solaris Synchronization (5.9.3)

- Solaris provides semaphores and conditions as we've already described
- *Adaptive mutexes* shift themselves from spinlocks to yielding if the kernel determines they're going to have to wait any real length of time for the lock
- *Turnstiles* are Solaris' method of queuing for adaptive mutexes
 - Turnstiles are dynamically allocated queues that implement priority inheritance
- Every synchronization mechanism available to the kernel is available to the user
 - ...except priority inheritance

LOOSE ENDS 3: BEYOND SEMAPHORES

Other Issues with Semaphores

- Mutexes and semaphores are useful, but they're low-level constructs
- Think of them as like `malloc()` and `free()` – they have all the same problems that can cause resource “leaks”
- So can we come up with the equivalent of garbage-collected objects?
- As a matter of fact...

NEXT TIME:
MONITORS, PIPES AND
SOCKETS
