

LECTURE 2.2

PAGING

COP4600

Dr. Matthew Gerber

3/2/2016

PAGING (8.5)


Paging (8.5.1)

- Take segmentation one step farther
- Break all of physical memory into fixed-sized blocks called *frames*
- Break all of logical memory into blocks of the same size, called *pages*
- When a process is to be executed, its *pages* are loaded into *frames*

Paged Addressing (8.5.1)

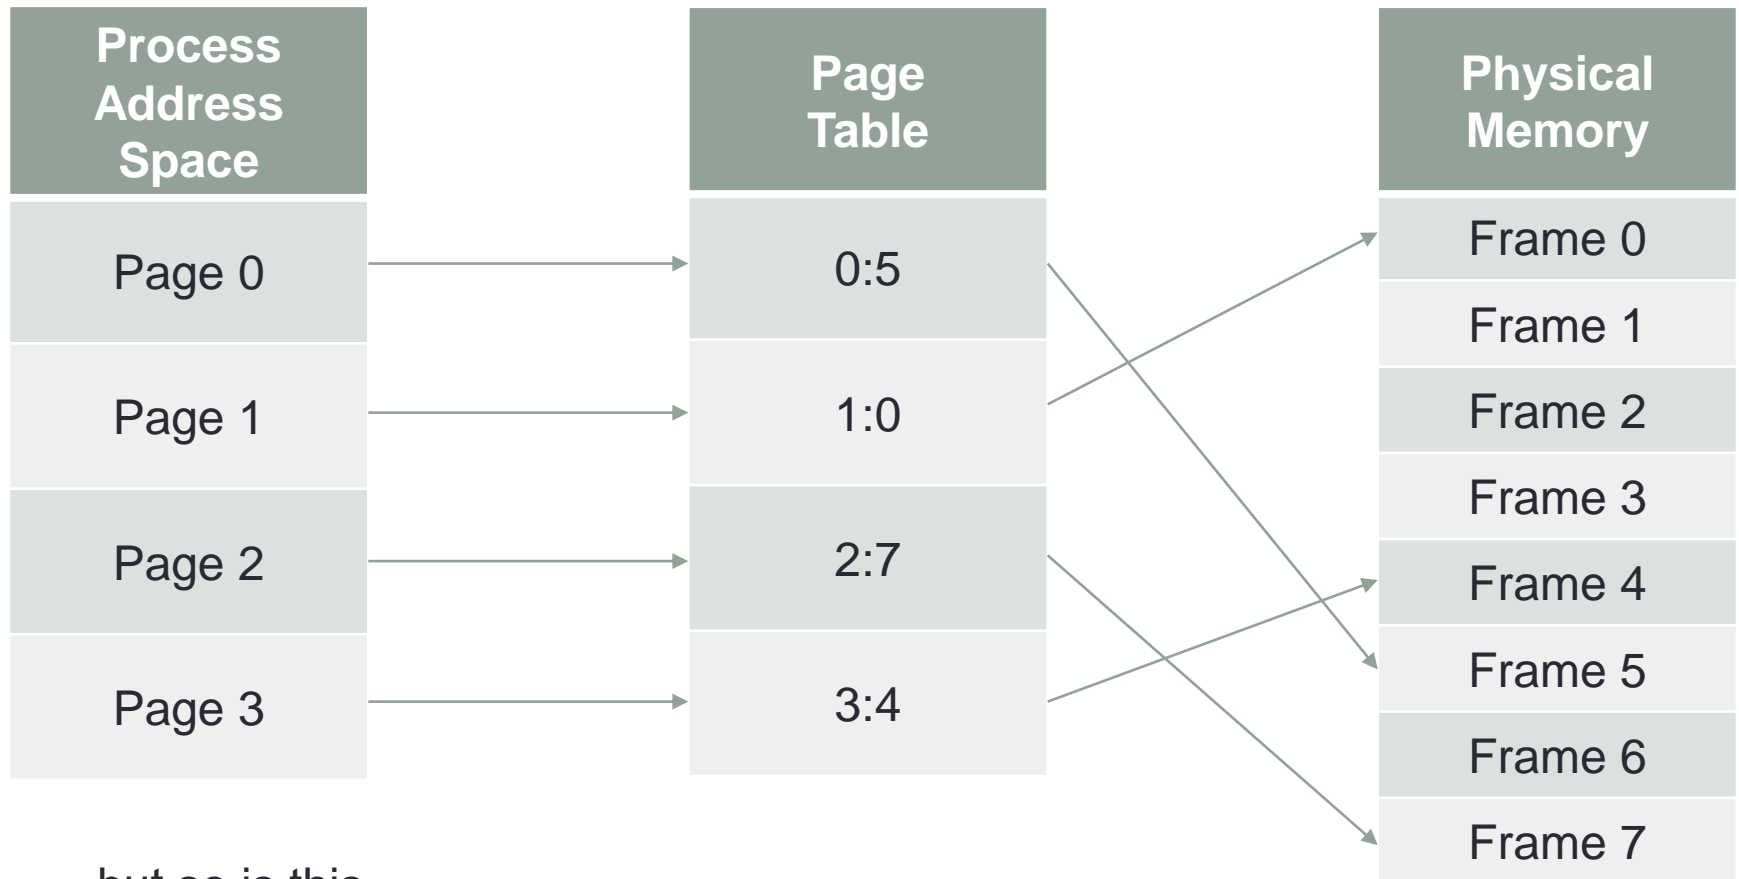
- Addresses again consist of two components: a *page number* and an *offset*
- The size of a page is invariably a power of 2, so that the address can be a single integer and we can extract the page number and offset using bitwise functions
- On every address access, translation is performed from the logical *page number* to the physical *frame*, and the offset is added to the result
- The translation is done by means of a (typically per-process) *page table*

The Page Table (8.5.1)




This is perfectly legal...

The Page Table (8.5.1)



The Page Table (8.5.1)



...but so is this

Addressing Again (8.5.1)

- This is dynamic relocation taken to its logical conclusion
- In the second example, assume our page size was 512 bytes
- Logical address 511 would map to the last byte of page 0, and be in frame 5
- Logical address 512 would map to the first byte of page 1, and be in frame 0
- *The program never has to care!*

Fragmentation and Page Size (8.5.1)

- Paging completely eliminates external fragmentation and the entire memory hole problem
- Paging does induce *internal* fragmentation
- We can easily observe that every process will waste an average of half the page size
- That's a good reason to have a small page size...
- ...but the management isn't free
- The good news is, since user-level programs simply don't care, we can let page sizes slowly grow over time without breaking anything

The Frame Table (8.5.1)

- As well as keeping track of pages per process, we also need to keep track of them system wide
- For every process, we have a page table
- For the system, we have a *frame table*
- The frame table has an entry for every physical frame that indicates:
 - Whether it is allocated, and if so
 - To which process, and
 - To which page of that process

Paging: Hardware Support (8.5.2)

- Even more so than simple base and limit registers, the hardware *must* support this
- Since we (usually) use a page table per process, context switch time is increased
- The page table(s) also have to be kept somewhere
- They're large enough that they need to be kept in main memory, but that can theoretically double memory access time
- Guess what happens next?

The Translation Look-Aside Buffer (8.5.2)

- A *translation look-aside buffer (TLB)* is a cache for page table entries
- With a TLB available, we only have to go to main memory to get a frame number when the page table entry isn't already in the TLB – this is a *TLB miss*
- TLBs may contain page table entries for only one process or may be able to contain entries for an arbitrary number of processes
- As with any cache, the TLB has a *hit ratio*

Paging: Protection (8.5.3)

- As well as having frame numbers, page table entries can simply be marked as not allocated
- If a process tries to access a non-allocated page, it incurs a protection violation
- A process's page table obviously does not always need to map the entire address space; if it doesn't, access to space that isn't in the page table can be presumed to be non-allocated
- We can also define specific pages as read-only, read-write, or even execute-only

Paging: Page Sharing (8.5.4)

- If more than one process uses identical pages for code (as opposed to data), and those code pages *never modify themselves*, then there's no reason those pages need to be unique for each process
- Useful for programs that need to be run more than once on servers
- Also useful for shared libraries
- The read-only property of shared code *must* be enforced by the OS

Hierarchical Page Tables (8.6.1)

- 32-bit addresses allow for a pretty big address space: 4,294,967,296 bytes
- With small pages, if we use a nontrivial amount of the space, the page table itself can reach into the megabytes
- One solution to this is to page the page table
- By the time we get to 64-bit addresses and the resulting 18,446,744,073,709,551,616-byte address space, that isn't enough to make the size manageable...
- ...but as it turns out, we aren't *actually* using 64-bit addresses now

Case Study: IA32 (8.7.1)

- Intel's legacy IA-32 architecture uses a two-level paging table for 4KB pages
- (It also still allows for old-school segmented addresses)
- IA-32 also supports a *page address extension* scheme that effectively expands paging by process
 - *Processes* are still limited to 4GB
 - The *system* can address up to 64GB
 - This *is* actually supported on some operating systems

Case Study: x86-64 (8.7.2)

- Intel's (really, AMD's) x86-64 architecture provides a four-level paging table, with page sizes of 4KB, 2MB, and 1GB
- Only 48 bits are actually used
- That's enough for 262,144GB, so it's still probably going to be pretty safe for a while
 - PAE is also supported if we're really worried
 - That brings the total system-addressable memory to 4,096TB

What About Swapping Pages?

- Remember that we advanced the idea of swapping out processes
- This isn't actually practical in most cases because of the sheer overhead involved in swapping them out or bringing them back in
- But what if we could swap out only certain *pages* of a process?
 - Initialization or rarely-used code fragments
 - Historical data that may not need to be referenced for a long time
 - And so on...

NEXT TIME: VIRTUAL
MEMORY
