

# LECTURE 2.5

## PAGE REPLACEMENT AND ACCESS TIME

---

COP4600

Dr. Matthew Gerber

3/21/2015

# PAGE REPLACEMENT ALGORITHMS

---

# Page Replacement

- In virtual memory, we constantly have the following case:
  - We need to swap a page from disk into main memory, and
  - There are no free frames remaining in main memory – whether in our program's allocated frames, or in main memory in its entirety
- In this case, we need to choose a *victim frame* to swap out first
- The choice of the victim frame is made by a *page replacement algorithm*

# FIFO Page Replacement (9.4.2)

- Replace the page that was loaded earliest
- So with three frames and references ordered as follows:

0, 1, 1, 1, 1, 0, 3, 0, 2

...frame 0 is replaced when frame 2 is referenced

- Simple to implement
- Poor performance to the point of often being outperformed by random replacement
- A particular issue – consider references ordered as follows:  
1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- How many page faults with four frames? How about with five?

# Belady's Anomaly and Stack Algorithms

- For some case with some page replacement algorithms, the number of page faults can *increase* when the number of available frames rises
- Algorithms with this problem are said to suffer from Belady's Anomaly
- Page replacement algorithms with this problem are generally considered to be poor algorithms
- Algorithms are called *stack algorithms* if the set of pages with fewer frames available is always a strict subset of the set of pages with more frames available
- Stack algorithms cannot suffer from Belady's anomaly

# Optimal Page Replacement (9.4.3)

- The optimal algorithm can be proven to be the best possible page replacement algorithm
  - (or at least a best possible page replacement algorithm, and that's what matters)
- Its definition is simple:

*Replace the page that will not be used for the longest period of time.*

- Advantage it is that it is in fact an optimal algorithm
- Disadvantage is that it cannot actually be implemented in the general case
  - If I can predict all memory accesses, what else can I predict?

# LRU Page Replacement (9.4.4)

- If we can't choose which page won't be used *in* the longest, we can instead replace that hasn't been used *for* the longest
- So with three frames and references ordered as follows:  
0, 1, 1, 1, 1, 0, 3, 0, 2  
...frame 1 is replaced when frame 2 is referenced
- Is a stack algorithm
- Generally has good performance
- Approximations are often used in practice

# LRU Implementations

- **Counter implementation:** Associate a CPU clock time with each page table entry
- **Stack implementation:** Maintain a list of all referenced pages in a doubly-linked list
- Either of these absolutely requires hardware to be effective
- In fact, even approximations of LRU require hardware



# LRU Approximations (9.4.5)

- **Second Chance**

- Each page in the table is given a *reference bit*
- The reference bit is set to 1 whenever the page is referenced
- Select a page for replacement as in FIFO, but if the reference bit of the selected page is 1, set its reference bit to 0 and pass over it
- Degenerates to FIFO if all bits are 1
- Can be enhanced to take into account page modification (Why?)

- **Additional Reference Bits**

- Like Second Chance, each page has a reference bit
- The reference bit is treated as the top bit of a reference *integer*
- Periodically, the reference integer is right-shifted
- When a page needs to be replaced, select a page with a minimum reference integer

# Counting Algorithms (9.4.6)

- Keep a counter of the number of references that have been made to each page
- **Least Frequently Used**
  - So with three frames and references ordered as follows:  
0, 1, 1, 1, 1, 0, 3, 0, 2  
...frame 3 is replaced when frame 2 is referenced
  - Expensive implementation, mediocre performance
  - Only useful for specific applications
- **Most Frequently Used**
  - So with three frames and references ordered as follows:  
0, 1, 1, 1, 1, 0, 3, 0, 2  
...frame 1 is replaced when frame 2 is referenced
  - Theory is that the page with the smallest count was just swapped in
  - Rarely actually useful

# MEMORY ACCESS TIME

---

# Memory Access Time with TLBs

- Recall that page tables are stored in memory
- *Translation lookaside buffers* are page table caches stored on the processor, and take less time to access
- We assume that TLBs take effectively no time to access themselves
- Memory access time for a TLB hit is the same as the underlying physical memory access time
- Memory access time for a TLB miss is *twice* the physical memory access time
- Like any other caches, TLBs have *hit rates* and *miss rates*
- So with physical memory access time  $m$  and TLB hit rate  $r$ , *effective* memory access time is  $mr + 2m(1 - r)$

# Memory Access with TLBs - Example

Physical Access Time	TLB Hit Rate	Effective Access Time	Paging Slowdown
100ns	50%		
100ns	80%		
100ns	90%		
100ns	99%		
100ns	99.9%		
100ns	99.99%		
100ns	99.999%		
100ns	99.9999%		
100ns	99.99999%		

# Memory Access with TLBs - Example

Physical Access Time	TLB Hit Rate	Effective Access Time	Paging Slowdown
100ns	50%	150ns	50.00000%
100ns	80%	120ns	20.00000%
100ns	90%	110ns	10.00000%
100ns	99%	101ns	1.00000%
100ns	99.9%	100.1ns	0.10000%
100ns	99.99%	100.01ns	0.01000%
100ns	99.999%	100.001ns	0.00100%
100ns	99.9999%	100.0001ns	0.00010%
100ns	99.99999%	100.00001ns	0.00001%

# Access Time with Virtual Memory

- Under virtual memory, when we need to access a page that isn't loaded into a frame, we have to go get it from disk
- For purposes of calculating the time involved, we can think of main memory as one big cache for virtual memory, with the miss rate being the page fault rate
- So with page fault rate  $f$ , memory access time  $m$  and disk swap time  $d$ , effective access time is  $(1 - f)m + fd$

# Access Time with Virtual Memory - Example

Mem. Acc. Time	Disk Swap Time	Page Fault Rate	Eff. Access Time	VM Slowdown
100ns	10,000,000ns	10.00000%		
100ns	10,000,000ns	1.00000%		
100ns	10,000,000ns	0.10000%		
100ns	10,000,000ns	0.01000%		
100ns	10,000,000ns	0.00100%		
100ns	10,000,000ns	0.00010%		
100ns	10,000,000ns	0.00001%		
100ns	200,000ns	10.00000%		
100ns	200,000ns	1.00000%		
100ns	200,000ns	0.10000%		
100ns	200,000ns	0.01000%		
100ns	200,000ns	0.00100%		
100ns	200,000ns	0.00010%		
100ns	200,000ns	0.00001%		



# Access Time with Virtual Memory - Example

Mem. Acc. Time	Disk Swap Time	Page Fault Rate	Eff. Access Time		VM Slowdown
100ns	10,000,000ns	10.000000%	1,000,090	ns	999990.000000%
100ns	10,000,000ns	1.000000%	100,099	ns	99999.000000%
100ns	10,000,000ns	0.100000%	10,099.9	ns	9999.900000%
100ns	10,000,000ns	0.010000%	1,099.99	ns	999.990000%
100ns	10,000,000ns	0.001000%	199.999	ns	99.999000%
100ns	10,000,000ns	0.000100%	109.9999	ns	9.999900%
100ns	10,000,000ns	0.000010%	100.99999	ns	0.999990%
100ns	200,000ns	10.000000%	20,090	ns	19990.000000%
100ns	200,000ns	1.000000%	2,099	ns	1999.000000%
100ns	200,000ns	0.100000%	299.9	ns	199.900000%
100ns	200,000ns	0.010000%	119.99	ns	19.990000%
100ns	200,000ns	0.001000%	101.999	ns	1.999000%
100ns	200,000ns	0.000100%	100.1999	ns	0.199900%
100ns	200,000ns	0.000010%	100.01999	ns	0.019990%

# Other Memory Access Time Factors

- We've considered:
  - The TLB
  - Virtual Memory
- Other factors include:
  - Caches
  - Addressing Modes
  - Specialized Operations
  - NUMA

# NEXT TIME: FRAME ALLOCATION AND EXAMPLES

---